3. If both policy engines approve, a Conclave Pod (CC) is spawned on each organization's OpenShift project.

CC-1

Spark Job — Read

Write

MPC Job — Read

4. Intermediate data is read & written from each organization's Swift volume.

Organization 1's Swift Volume

Organization 1 Policy Engine

1. User submits workflow

Conclave-web

2. Workflow is verified by each organization's policy engine.

* Pod communication is necessary for MPC jobs.

User's Swift volume

Organization 2 Policy Engine

CC-2

Spark Job — Read

Write

MPC Job — Read

Organization 2's Swift Volume

5. When the computation is finished, output is written to a Swift volume that the User can access